

Intro to Data Science - HW 6

Copyright Jeffrey Stanton 2022, Jeffrey Saltz, and Jasmina Tacheva

```
# Enter your name here: Lily Coyle
```

Attribution statement: (choose only one and delete the rest)

```
# 1. I did this homework by myself, with help from the book and the professor.
```

This module: Data visualization is important because many people can make sense of data more easily when it is presented in graphic form. As a data scientist, you will have to present complex data to decision makers in a form that makes the data interpretable for them. From your experience with Excel and other tools, you know that there are a variety of **common data visualizations** (e.g., pie charts). How many of them can you name?

The most powerful tool for data visualization in R is called **ggplot**. Written by computer/data scientist **Hadley Wickham**, this **** graphics grammar **** tool builds visualizations in layers. This method provides immense flexibility, but takes a bit of practice to master.

Step 1: Make a copy of the data

- A. Read the **New York State COVID Testing** dataset we used in HW 3 & 4 from this URL: https://data-science-intro.s3.us-east-2.amazonaws.com/NYS_COVID_Testing_.csv (https://data-science-intro.s3.us-east-2.amazonaws.com/NYS_COVID_Testing_.csv) into a new dataframe called **df**.

```
library(tidyverse)
```

```
## — Attaching packages ————— tidyverse 1.3.1 —
```

```
## ✓ ggplot2 3.3.5      ✓ purrr   0.3.4
## ✓ tibble  3.1.6      ✓ dplyr   1.0.7
## ✓ tidyr   1.2.0      ✓ stringr 1.4.0
## ✓ readr   2.1.2      ✓ forcats 0.5.1
```

```
## — Conflicts ————— tidyverse_conflicts() —
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
```

```
df <- read_csv("https://data-science-intro.s3.us-east-2.amazonaws.com/NYS_COVID_Testing_.csv")
```

```
## Rows: 7383 Columns: 7
```

```
## — Column specification —————  
## Delimiter: ","  
## chr (3): TestDate, AgeGroup, AgeCategory  
## dbl (4): PositiveCases, TotalTests, Year, id
```

```
##  
## i Use `spec()` to retrieve the full column specification for this data.  
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

B. Your dataframe, **df**, contains a so-called **multivariate time series**: a sequence of measurements on COVID tests and results captured repeatedly over time (March 2020 - January 2022). Familiarize yourself with the nature of the time variable **TestDate**.

How often were these measurements taken (in other words, at what frequency were the variables measured)? Put your answer in a comment.

```
#The measurements were taken three times a day. Each age group was taken once a day.
```

C. What is the data type of **TestDate**? Explain in a comment.

```
#The data type of TestDate is character.
```

D. To properly display the **TestDate** values as dates in our plots, we need to convert **TestDate** to date format with the **as.Date()** function. Run the code below and check the data type of the variable again to make sure it is not coded as text anymore:

```
df$TestDate<-as.Date(df$TestDate, format = "%m/%d/%Y")
```

```
str(df)
```

```
## spec_tbl_df [7,383 × 7] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
## $ TestDate      : Date[1:7383], format: "2020-03-04" "2020-03-04" ...
## $ AgeGroup      : chr [1:7383] "1 to 4" "5 to 19" "< 1" "1 to 4" ...
## $ PositiveCases : num [1:7383] 0 0 0 0 6 0 1 0 8 2 ...
## $ TotalTests    : num [1:7383] 1 3 1 7 17 2 10 3 52 18 ...
## $ AgeCategory   : chr [1:7383] "children" "children" "children" "children" ...
## $ Year          : num [1:7383] 2020 2020 2020 2020 2020 2020 2020 2020 2020 2020
## ...
## $ id            : num [1:7383] 8 12 16 17 22 27 32 37 42 47 ...
## - attr(*, "spec")=
## .. cols(
## ..   TestDate = col_character(),
## ..   AgeGroup = col_character(),
## ..   PositiveCases = col_double(),
## ..   TotalTests = col_double(),
## ..   AgeCategory = col_character(),
## ..   Year = col_double(),
## ..   id = col_double()
## .. )
## - attr(*, "problems")=<externalptr>
```

Step 2: Clean up the NAs and create subsets

- A. It is always good practice, when you first start working with a dataset, to explore it for missing values. Check the **TotalTests** and **PositiveCases** for missing values. Are there any? What does empty output suggest about the number of missing observations?

Hint: use `is.na()`

```
df[is.na(df$TotalTests),] #there are 22 rows with missing observations
```

```
## # A tibble: 22 × 7
##   TestDate   AgeGroup PositiveCases TotalTests AgeCategory      Year    id
##   <date>     <chr>         <dbl>     <dbl> <chr>         <dbl> <dbl>
## 1 2020-04-22 1 to 4           24         NA children    2020   542
## 2 2020-04-22 5 to 19        289         NA children    2020   547
## 3 2020-04-22 < 1           4         NA children    2020   552
## 4 2021-12-17 1 to 4          533         NA children    2021  7186
## 5 2021-12-17 5 to 19       3986         NA children    2021  7191
## 6 2021-12-17 < 1          146         NA children    2021  7196
## 7 2020-04-22 35 to 44        992         NA middle-aged_adults 2020   545
## 8 2020-04-22 45 to 54       1089         NA middle-aged_adults 2020   546
## 9 2021-12-17 35 to 44       3330         NA middle-aged_adults 2021  7189
## 10 2021-12-17 45 to 54      2138         NA middle-aged_adults 2021  7190
## # ... with 12 more rows
```

```
df[is.na(df$PositiveCases),] #there are 11 rows with 0 positive cases
```

```
## # A tibble: 11 × 7
##   TestDate   AgeGroup PositiveCases TotalTests AgeCategory   Year   id
##   <date>     <chr>         <dbl>     <dbl> <chr>         <dbl> <dbl>
## 1 2021-12-15 1 to 4           NA       9694 children      2021  7164
## 2 2021-12-15 5 to 19          NA      54492 children      2021  7169
## 3 2021-12-15 < 1            NA       1430 children      2021  7174
## 4 2021-12-15 35 to 44         NA      40170 middle-aged_adults 2021  7167
## 5 2021-12-15 45 to 54         NA      32945 middle-aged_adults 2021  7168
## 6 2021-12-15 55 to 64         NA      29690 older_adults      2021  7170
## 7 2021-12-15 65 to 74         NA      16472 older_adults      2021  7171
## 8 2021-12-15 75 to 84         NA       8168 older_adults      2021  7172
## 9 2021-12-15 85 +            NA       6055 older_adults      2021  7173
## 10 2021-12-15 20 to 24        NA      24679 young_adults      2021  7165
## 11 2021-12-15 25 to 34        NA      54164 young_adults      2021  7166
```

B. There is an R package called **imputeTS** specifically designed to repair missing values in time series data. We will use this instead of the simpler way, **mean substitution**, because it tends to be more accurate. The **na_interpolation()** function in this package takes advantage of a unique characteristic of time series data: neighboring points in time can be used to guess about a missing value in between. Use this function on each of the two numeric variables in **df** and don't forget to **update** them by overwriting them with the output of the **na_interpolation()** function.

```
library("imputeTS")
```

```
## Registered S3 method overwritten by 'quantmod':
##   method           from
##   as.zoo.data.frame zoo
```

```
df$TotalTests <- na_interpolation(df$TotalTests)
df$PositiveCases <- na_interpolation(df$PositiveCases)
```

C. Run the code from A to make sure there is no more missing data:

```
df[is.na(df$TotalTests),]
```

```
## # A tibble: 0 × 7
## # ... with 7 variables: TestDate <date>, AgeGroup <chr>, PositiveCases <dbl>,
## #   TotalTests <dbl>, AgeCategory <chr>, Year <dbl>, id <dbl>
```

```
df[is.na(df$PositiveCases),]
```

```
## # A tibble: 0 × 7
## # ... with 7 variables: TestDate <date>, AgeGroup <chr>, PositiveCases <dbl>,
## #   TotalTests <dbl>, AgeCategory <chr>, Year <dbl>, id <dbl>
```

D. As we've done before, let's create a new variable which is the ratio of **PositiveCases** to **TotalTests** - save it

as an additional variable in **df** called **PositivityRate**:

```
df$PositivityRate <- df$PositiveCases / df$TotalTests
```

E. Create a subset of **df** containing **only the records for children**. Save it in a new dataframe called **dfChildren**. Make sure this new df has **2,010 observations and 6 variables**.

```
dfChildren <- df[df$AgeCategory == "children",]
```

F. Create a subset of **df** containing only the records for **young adults**. Save it in a new dataframe called **dfYA**.

```
dfYA <- df[df$AgeCategory == "young_adults",]
```

G. Using the same logic, create 2 more subsets of **df**: one containing only the records for **middle-aged adults** (call it **dfMA**), and another one with only the data of **older adults** - **dfOA**. After this step, you should have a total of 4 subsets: - dfChildren - dfYA - dfMA - dfOA

```
dfMA <- df[df$AgeCategory == "middle-aged_adults",]  
dfOA <- df[df$AgeCategory == "older_adults",]
```

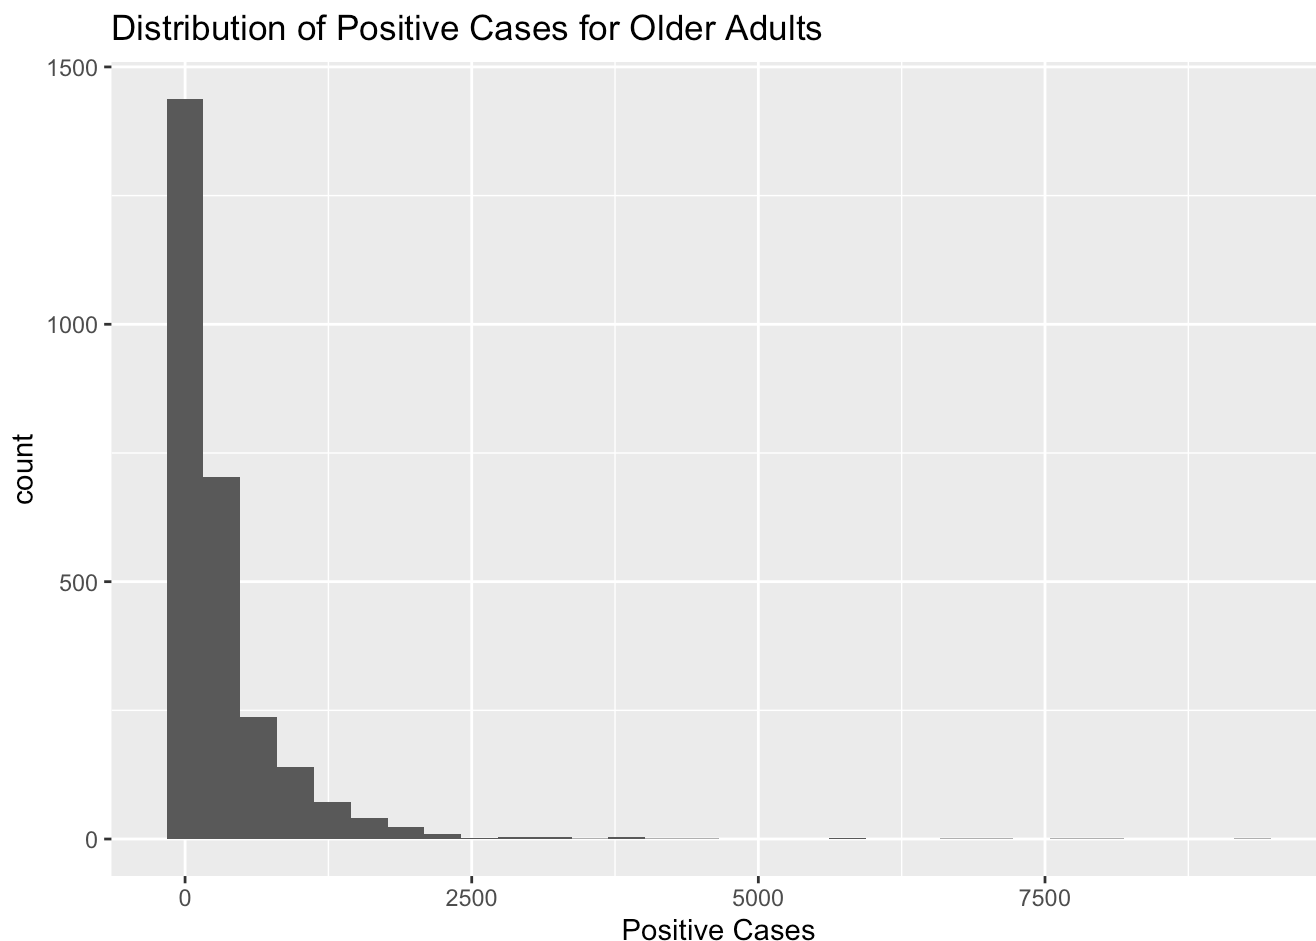
Step 3: Use ggplot to explore the distribution of each variable

Don't forget to install and library the ggplot2 package. Then:

A. Create a histogram for **PositiveCases** in the **dfOA** dataframe. Be sure to add a title and briefly describe what the histogram means in a comment.

```
hist0A <- ggplot(dfOA, aes(x=PositiveCases)) + geom_histogram() + ylab("count") + xlab("Positive Cases") + ggtitle("Distribution of Positive Cases for Older Adults")  
hist0A
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



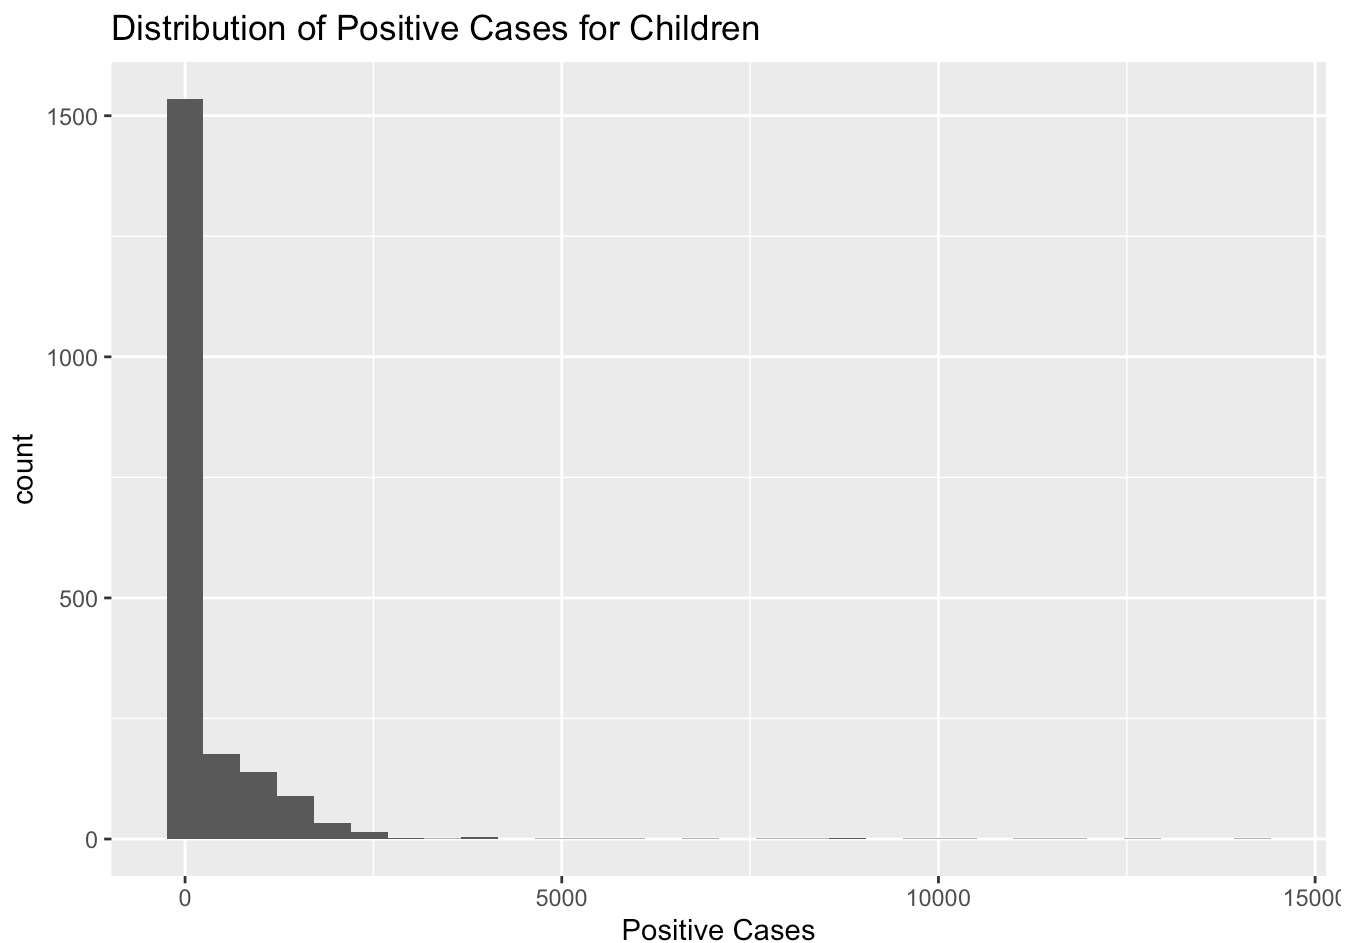
#The histogram means that the highest distribution of positive cases came at the start

B. Create histograms (using **ggplot**) of the **PositiveCases** variable in each of the other three subsets from Step 2G.

For each histogram, comment on its shape - what information can we glean from it?

```
histChildren <- ggplot(dfChildren, aes(x=PositiveCases)) + geom_histogram() + ylab("count") + xlab("Positive Cases") + ggtitle("Distribution of Positive Cases for Children")
histChildren
```


```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



#Once again the highest distribution of positive cases came at the start and then began to taper off and stabilize

```
histMA <- ggplot(dfMA, aes(x=PositiveCases)) + geom_histogram() + ylab("count") + xlab("Positive Cases") + ggtitle("Distribution of Positive Cases for Middle Aged Adults")
histMA
```

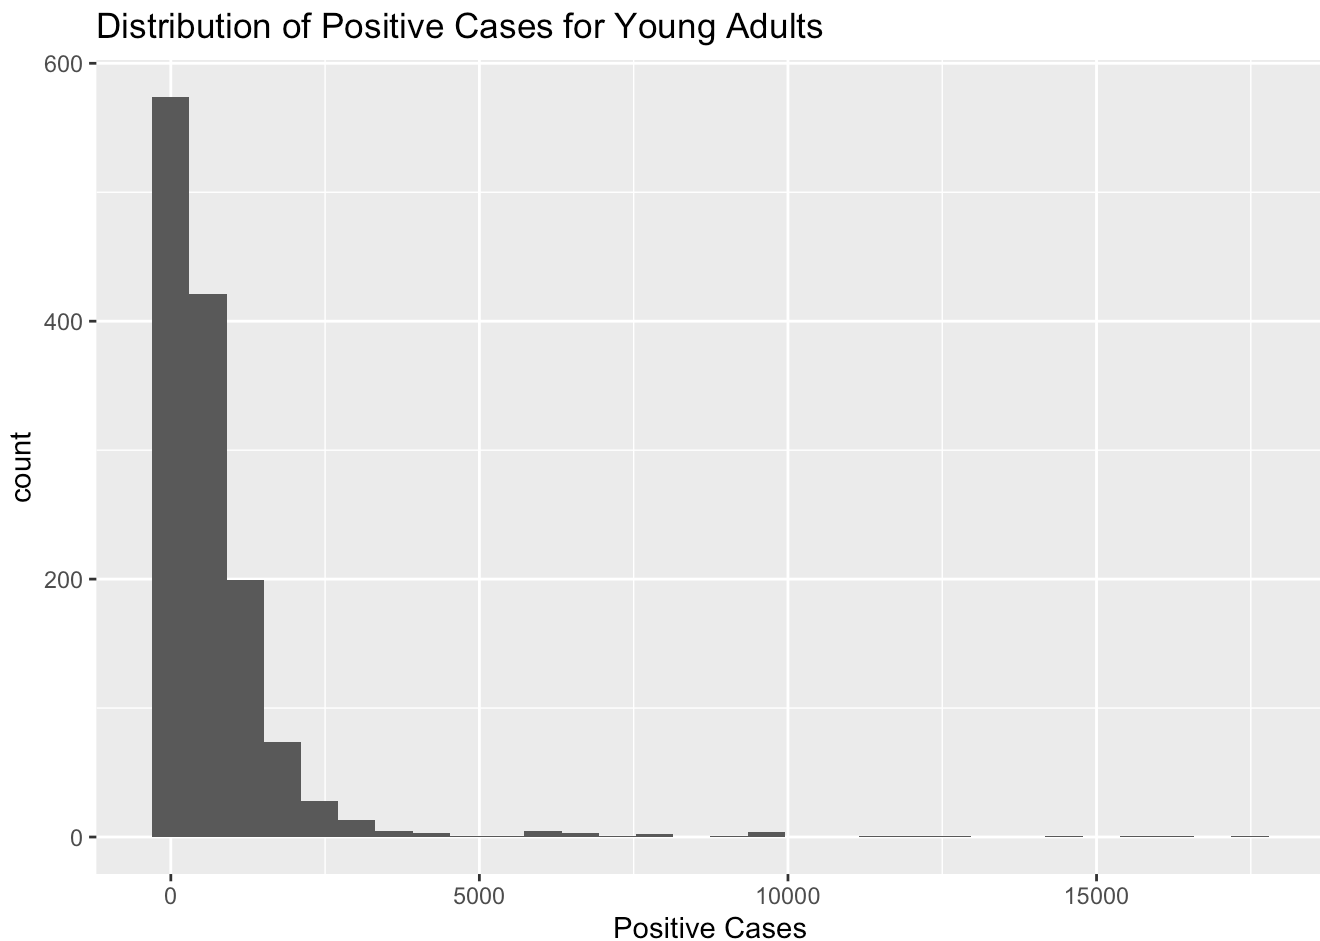
```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```



#In this case the spike was also at the start but it took longer for it to go down

```
histYA <- ggplot(dfYA, aes(x=PositiveCases)) + geom_histogram() + ylab("count") + xlab("Positive Cases") + ggtitle("Distribution of Positive Cases for Young Adults")
histYA
```

```
## `stat_bin()` using `bins = 30`. Pick better value with `binwidth`.
```


#Similar to the one above this one had a large spike right at the beginning and tapered off at a slower pace

Step 4: Explore how the data changes over time

A. These data were collected in a period of almost 2 years. You can thus observe changes over time with the help of a line chart. Let's focus on the **dfOA** subset first:


Create a **line chart**, with **TestDate** on the X-axis and **PositivityRate** on the Y-axis.

```
line0A <- ggplot(dfOA, aes(x= TestDate, y = PositivityRate)) + geom_line(aes(y=PositivityRate, color = "green"))
line0A
```




B. Next, create similar graphs for each of the other three subsets. Change the **color** of the line plots (any color you want).


```
lineMA <- ggplot(dfMA, aes(x= TestDate, y = PositivityRate)) + geom_line(aes(y=PositivityRate, color = "red"))
lineMA
```



```
lineYA <- ggplot(dfYA, aes(x= TestDate, y = PositivityRate)) + geom_line(aes(y=PositivityRate, color = "red"))
lineYA
```



```
lineChildren <- ggplot(dfChildren, aes(x= TestDate, y = PositivityRate)) + geom_line  
(aes(y=PositivityRate, color = "red"))  
lineChildren
```




C. In a comment, talk about the insights you got from the line charts you created - can you spot any trends within and between the line charts?

#The biggest spike across all age groups came at the start of the pandemic and didn't rise dramatically again until the start of omicron


D. Finally, using the **dfOA** subset, create a **scatter plot**, showing **TotalTests** on the x axis, **PositiveCases** on the y axis, and having the **color and size** of the point represent **Year**.

```
scatter0A <- ggplot(dfOA, aes(x = TotalTests, y = PositiveCases, col = Year, size = Year)) + geom_point()
scatter0A
```



E. Create a similar scatter plot for the **dfYA** subset.

```
scatterYA <- ggplot(dfYA, aes(x = TotalTests, y = PositiveCases, col = Year, size = Year)) + geom_point()
scatterYA
```



F. Interpret these visualizations what insight do they provide?

#These visualizations show the number of positive covid tests in relation to the year in which they were taken in. In some instances, closest to 2021, the cases were the highest. However, it is evident the highest number of positive cases was in 2020 and in 2022 they are getting lower.